

Data Analysis And Practical - 4

Shad Jamil

CSC/21/45

In [4]:

```
1 import numpy as np
2 import pandas as pd
3 #Creating a dataframe with 3 columns and 50 rows
4 dat = pd.DataFrame(np.random.rand(50,3), columns=['Col1','Col2','Col3'])
5 #Replacing 10% values with null values
6 x = int(0.1*dat.size)
7 indices_to_replace = np.unravel_index(np.random.choice(dat.size, x, replace=True), dat.shape)
8 dat.iloc[indices_to_replace] = np.nan
9 #Printing Results
10 print('Acquired DataFrame:\n')
11 print(dat)
```

Acquired DataFrame:

	Col1	Col2	Col3
0	0.682881	0.206784	0.013455
1	0.783925	0.393434	0.049236
2	0.772594	0.206481	0.288398
3	0.492280	0.484211	0.194326
4	0.197248	0.068652	0.268006
5	0.892160	0.080545	0.433682
6	0.305926	0.422013	0.604318
7	0.313714	0.043631	0.774311
8	0.339754	0.965632	0.034533
9	0.148361	0.725745	0.065957
10	0.454817	0.639184	0.845437
11	0.846323	0.828858	0.199262
12	0.021332	0.195676	0.139573
13	0.922980	0.738134	0.545251
14	0.967218	0.636276	0.668695
15	NaN	NaN	NaN
16	NaN	NaN	NaN
17	0.669454	0.831103	0.313997
18	0.185592	0.747234	0.676615
19	0.541976	0.502911	0.676689
20	0.280902	0.775334	0.883078
21	NaN	NaN	NaN
22	0.419839	0.531416	0.274781
23	0.683998	0.927282	0.632482
24	NaN	NaN	NaN
25	NaN	NaN	NaN
26	0.817794	0.678031	0.030699
27	0.356871	0.192558	0.062831
28	0.487123	0.912645	0.873654
29	0.048407	0.658765	0.620407
30	0.494716	0.308894	0.938305
31	0.820723	0.104469	0.264013
32	0.918657	0.797679	0.222038
33	NaN	NaN	NaN
34	NaN	NaN	NaN
35	0.900773	0.481834	0.259489
36	0.529321	0.710574	0.490034
37	NaN	NaN	NaN
38	NaN	NaN	NaN
39	0.458132	0.604310	0.754716
40	0.083181	0.165589	0.275054
41	0.396066	0.545951	0.461533
42	NaN	NaN	NaN
43	0.663116	0.556754	0.263829
44	0.030904	0.587509	0.907292
45	0.381704	0.326459	0.607293
46	NaN	NaN	NaN
47	0.010397	0.444512	0.551197
48	0.214805	0.222936	0.274977
49	NaN	NaN	NaN

```
In [5]: 1 #(a) - Identify and Count missing values
        2 #Identifying missing values
        3 missing_values_identified_df = dat.isnull()
        4 #Counting missing values
        5 total_missing_values = missing_values_identified_df.sum().sum()
        6 #Printing Results
        7 print('DataFrame With Missing Values:\n')
        8 print(missing_values_identified_df)
        9 print('Total Missing Values:', total_missing_values)
```

DataFrame With Missing Values:

	Col1	Col2	Col3
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
5	False	False	False
6	False	False	False
7	False	False	False
8	False	False	False
9	False	False	False
10	False	False	False
11	False	False	False
12	False	False	False
13	False	False	False
14	False	False	False
15	True	True	True
16	True	True	True
17	False	False	False
18	False	False	False
19	False	False	False
20	False	False	False
21	True	True	True
22	False	False	False
23	False	False	False
24	True	True	True
25	True	True	True
26	False	False	False
27	False	False	False
28	False	False	False
29	False	False	False
30	False	False	False
31	False	False	False
32	False	False	False
33	True	True	True
34	True	True	True
35	False	False	False
36	False	False	False
37	True	True	True
38	True	True	True
39	False	False	False
40	False	False	False
41	False	False	False
42	True	True	True
43	False	False	False
44	False	False	False
45	False	False	False
46	True	True	True
47	False	False	False
48	False	False	False
49	True	True	True

Total Missing Values: 36

```
In [6]: 1  #(b) - Dropping column with more than 5 null values  
2  #Checking for columns with more than 5 null values  
3 columns_to_drop = dat.columns[dat.isnull().sum()>5]  
4  #Dropping appropriate columns  
5 columns_dropped_df = dat.drop(columns=columns_to_drop)  
6  #Printing Results  
7 print('DataFrame With Columns Dropped:\n')  
8 print(columns_dropped_df)  
9
```

DataFrame With Columns Dropped:

Empty DataFrame

Columns: []

Index: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49]

```
In [7]: 1 #(c) - Identify the row label having maximum of the sum of all values in c
2 #Identifying the row label having max sum
3 max_sum_row = dat.sum(axis=1).idxmax()
4 #Dropping appropriate column
5 max_sum_row_dropped_df = dat.drop(max_sum_row)
6 #Printing Results
7 print('Row Index with Max Sum:', max_sum_row, '\n')
8 print('DataFrame with Row With Max Sum Dropped:\n')
9 print(max_sum_row_dropped_df)
```

Row Index with Max Sum: 28

DataFrame with Row With Max Sum Dropped:

	Col1	Col2	Col3
0	0.682881	0.206784	0.013455
1	0.783925	0.393434	0.049236
2	0.772594	0.206481	0.288398
3	0.492280	0.484211	0.194326
4	0.197248	0.068652	0.268006
5	0.892160	0.080545	0.433682
6	0.305926	0.422013	0.604318
7	0.313714	0.043631	0.774311
8	0.339754	0.965632	0.034533
9	0.148361	0.725745	0.065957
10	0.454817	0.639184	0.845437
11	0.846323	0.828858	0.199262
12	0.021332	0.195676	0.139573
13	0.922980	0.738134	0.545251
14	0.967218	0.636276	0.668695
15	NaN	NaN	NaN
16	NaN	NaN	NaN
17	0.669454	0.831103	0.313997
18	0.185592	0.747234	0.676615
19	0.541976	0.502911	0.676689
20	0.280902	0.775334	0.883078
21	NaN	NaN	NaN
22	0.419839	0.531416	0.274781
23	0.683998	0.927282	0.632482
24	NaN	NaN	NaN
25	NaN	NaN	NaN
26	0.817794	0.678031	0.030699
27	0.356871	0.192558	0.062831
29	0.048407	0.658765	0.620407
30	0.494716	0.308894	0.938305
31	0.820723	0.104469	0.264013
32	0.918657	0.797679	0.222038
33	NaN	NaN	NaN
34	NaN	NaN	NaN
35	0.900773	0.481834	0.259489
36	0.529321	0.710574	0.490034
37	NaN	NaN	NaN
38	NaN	NaN	NaN
39	0.458132	0.604310	0.754716
40	0.083181	0.165589	0.275054
41	0.396066	0.545951	0.461533
42	NaN	NaN	NaN
43	0.663116	0.556754	0.263829
44	0.030904	0.587509	0.907292
45	0.381704	0.326459	0.607293
46	NaN	NaN	NaN
47	0.010397	0.444512	0.551197
48	0.214805	0.222936	0.274977
49	NaN	NaN	NaN

In [8]:

```
1  #(d) - Sort the dataframe on the basis of the first column
2 sorted_df = dat.sort_values(by='Col1')
3  #Printing Results
4 print('DataFrame Sorted by Col1:\n')
5 print(sorted_df)
6
```


DataFrame Sorted by Col1:

	Col1	Col2	Col3
47	0.010397	0.444512	0.551197
12	0.021332	0.195676	0.139573
44	0.030904	0.587509	0.907292
29	0.048407	0.658765	0.620407
40	0.083181	0.165589	0.275054
9	0.148361	0.725745	0.065957
18	0.185592	0.747234	0.676615
4	0.197248	0.068652	0.268006
48	0.214805	0.222936	0.274977
20	0.280902	0.775334	0.883078
6	0.305926	0.422013	0.604318
7	0.313714	0.043631	0.774311
8	0.339754	0.965632	0.034533
27	0.356871	0.192558	0.062831
45	0.381704	0.326459	0.607293
41	0.396066	0.545951	0.461533
22	0.419839	0.531416	0.274781
10	0.454817	0.639184	0.845437
39	0.458132	0.604310	0.754716
28	0.487123	0.912645	0.873654
3	0.492280	0.484211	0.194326
30	0.494716	0.308894	0.938305
36	0.529321	0.710574	0.490034
19	0.541976	0.502911	0.676689
43	0.663116	0.556754	0.263829
17	0.669454	0.831103	0.313997
0	0.682881	0.206784	0.013455
23	0.683998	0.927282	0.632482
2	0.772594	0.206481	0.288398
1	0.783925	0.393434	0.049236
26	0.817794	0.678031	0.030699
31	0.820723	0.104469	0.264013
11	0.846323	0.828858	0.199262
5	0.892160	0.080545	0.433682
35	0.900773	0.481834	0.259489
32	0.918657	0.797679	0.222038
13	0.922980	0.738134	0.545251
14	0.967218	0.636276	0.668695
15	NaN	NaN	NaN
16	NaN	NaN	NaN
21	NaN	NaN	NaN
24	NaN	NaN	NaN
25	NaN	NaN	NaN
33	NaN	NaN	NaN
34	NaN	NaN	NaN
37	NaN	NaN	NaN
38	NaN	NaN	NaN
42	NaN	NaN	NaN
46	NaN	NaN	NaN
49	NaN	NaN	NaN

```
In [9]: 1 #(e) - Remove all duplicates from the first column
2         duplicates_dropped_df = dat.drop_duplicates(subset='Col1')
3         #Printing Results
4         print('DataFrame with Duplicates Removed from Col1:\n')
5         print(duplicates_dropped_df)
```

DataFrame with Duplicates Removed from Col1:

	Col1	Col2	Col3
0	0.682881	0.206784	0.013455
1	0.783925	0.393434	0.049236
2	0.772594	0.206481	0.288398
3	0.492280	0.484211	0.194326
4	0.197248	0.068652	0.268006
5	0.892160	0.080545	0.433682
6	0.305926	0.422013	0.604318
7	0.313714	0.043631	0.774311
8	0.339754	0.965632	0.034533
9	0.148361	0.725745	0.065957
10	0.454817	0.639184	0.845437
11	0.846323	0.828858	0.199262
12	0.021332	0.195676	0.139573
13	0.922980	0.738134	0.545251
14	0.967218	0.636276	0.668695
15	NaN	NaN	NaN
17	0.669454	0.831103	0.313997
18	0.185592	0.747234	0.676615
19	0.541976	0.502911	0.676689
20	0.280902	0.775334	0.883078
22	0.419839	0.531416	0.274781
23	0.683998	0.927282	0.632482
26	0.817794	0.678031	0.030699
27	0.356871	0.192558	0.062831
28	0.487123	0.912645	0.873654
29	0.048407	0.658765	0.620407
30	0.494716	0.308894	0.938305
31	0.820723	0.104469	0.264013
32	0.918657	0.797679	0.222038
35	0.900773	0.481834	0.259489
36	0.529321	0.710574	0.490034
39	0.458132	0.604310	0.754716
40	0.083181	0.165589	0.275054
41	0.396066	0.545951	0.461533
43	0.663116	0.556754	0.263829
44	0.030904	0.587509	0.907292
45	0.381704	0.326459	0.607293
47	0.010397	0.444512	0.551197
48	0.214805	0.222936	0.274977

```
In [11]: 1  #(f) - Find the correlation between first and second column and covariance  
2  #Correlation (1st - 2nd col)  
3 corr_1_2 = dat['Col1'].corr(dat['Col2'])  
4  #Covariance (2nd - 3rd col)  
5 cov_2_3 = dat['Col2'].cov(dat['Col3'])  
6  #Printing Results  
7 print('Correlation b/w 1st and 2nd Column:', corr_1_2)  
8 print('Covariance b/w 2nd and 3rd Column:', cov_2_3)
```

Correlation b/w 1st and 2nd Column: 0.1420138739574065

Covariance b/w 2nd and 3rd Column: 0.01353966590974579

```
In [12]: 1  #(g) - Detect the outliers and remove the rows having outliers
2  #Function to Detect and Remove Outliers
3 def remove_outliers(dat, zscore_thresh=3):
4     z_scores = np.abs((dat - dat.mean()) / dat.std())
5     outlier_mask = (z_scores < zscore_thresh).all(axis=1)
6     new_df = dat[outlier_mask]
7     return new_df
8  #Removing Outliers From Data
9 outliers_removed_df = remove_outliers(dat)
10  #Printing Results
11 print('DataFrame with Outliers Removed:\n')
12 print(outliers_removed_df)
13
```

DataFrame with Outliers Removed:

	Col1	Col2	Col3
0	0.682881	0.206784	0.013455
1	0.783925	0.393434	0.049236
2	0.772594	0.206481	0.288398
3	0.492280	0.484211	0.194326
4	0.197248	0.068652	0.268006
5	0.892160	0.080545	0.433682
6	0.305926	0.422013	0.604318
7	0.313714	0.043631	0.774311
8	0.339754	0.965632	0.034533
9	0.148361	0.725745	0.065957
10	0.454817	0.639184	0.845437
11	0.846323	0.828858	0.199262
12	0.021332	0.195676	0.139573
13	0.922980	0.738134	0.545251
14	0.967218	0.636276	0.668695
17	0.669454	0.831103	0.313997
18	0.185592	0.747234	0.676615
19	0.541976	0.502911	0.676689
20	0.280902	0.775334	0.883078
22	0.419839	0.531416	0.274781
23	0.683998	0.927282	0.632482
26	0.817794	0.678031	0.030699
27	0.356871	0.192558	0.062831
28	0.487123	0.912645	0.873654
29	0.048407	0.658765	0.620407
30	0.494716	0.308894	0.938305
31	0.820723	0.104469	0.264013
32	0.918657	0.797679	0.222038
35	0.900773	0.481834	0.259489
36	0.529321	0.710574	0.490034
39	0.458132	0.604310	0.754716
40	0.083181	0.165589	0.275054
41	0.396066	0.545951	0.461533
43	0.663116	0.556754	0.263829
44	0.030904	0.587509	0.907292
45	0.381704	0.326459	0.607293
47	0.010397	0.444512	0.551197
48	0.214805	0.222936	0.274977

```
In [13]: 1  #(h) - Discretize second column and create 5 bins  
2  #Filling Missing (NaN) Values in Col2 because pandas.cut cannot appropriate  
3 dat['Col2'] = dat['Col2'].fillna(dat['Col2'].mean())  
4  #Invoking pandas.cut function  
5 dat['Col2_Discretized'] = pd.cut(dat.iloc[:,1], bins=5, labels=False)  
6  #Printing Results  
7 print('DataFrame with Col2 Discretized into 5 bins:\n')  
8 print(dat)  
9
```

DataFrame with Col2 Discretized into 5 bins:

	Col1	Col2	Col3	Col2_Discretized
0	0.682881	0.206784	0.013455	0
1	0.783925	0.393434	0.049236	1
2	0.772594	0.206481	0.288398	0
3	0.492280	0.484211	0.194326	2
4	0.197248	0.068652	0.268006	0
5	0.892160	0.080545	0.433682	0
6	0.305926	0.422013	0.604318	2
7	0.313714	0.043631	0.774311	0
8	0.339754	0.965632	0.034533	4
9	0.148361	0.725745	0.065957	3
10	0.454817	0.639184	0.845437	3
11	0.846323	0.828858	0.199262	4
12	0.021332	0.195676	0.139573	0
13	0.922980	0.738134	0.545251	3
14	0.967218	0.636276	0.668695	3
15	NaN	0.506579	NaN	2
16	NaN	0.506579	NaN	2
17	0.669454	0.831103	0.313997	4
18	0.185592	0.747234	0.676615	3
19	0.541976	0.502911	0.676689	2
20	0.280902	0.775334	0.883078	3
21	NaN	0.506579	NaN	2
22	0.419839	0.531416	0.274781	2
23	0.683998	0.927282	0.632482	4
24	NaN	0.506579	NaN	2
25	NaN	0.506579	NaN	2
26	0.817794	0.678031	0.030699	3
27	0.356871	0.192558	0.062831	0
28	0.487123	0.912645	0.873654	4
29	0.048407	0.658765	0.620407	3
30	0.494716	0.308894	0.938305	1
31	0.820723	0.104469	0.264013	0
32	0.918657	0.797679	0.222038	4
33	NaN	0.506579	NaN	2
34	NaN	0.506579	NaN	2
35	0.900773	0.481834	0.259489	2
36	0.529321	0.710574	0.490034	3
37	NaN	0.506579	NaN	2
38	NaN	0.506579	NaN	2
39	0.458132	0.604310	0.754716	3
40	0.083181	0.165589	0.275054	0
41	0.396066	0.545951	0.461533	2
42	NaN	0.506579	NaN	2
43	0.663116	0.556754	0.263829	2
44	0.030904	0.587509	0.907292	2
45	0.381704	0.326459	0.607293	1
46	NaN	0.506579	NaN	2
47	0.010397	0.444512	0.551197	2
48	0.214805	0.222936	0.274977	0
49	NaN	0.506579	NaN	2